

Relational Algebra Operations

Fundamental operations:

- ☆ Select (σ)
- ☆ Project (Π)
- ☆ Union (\cup)
- ☆ Set Difference ($-$)
- ☆ Cartesian Product (\times)
- ☆ Rename (ρ)

Relational Algebra Operations

Additional operations:

- ☆ Set intersection (\cap)
- ☆ Assignment operation ($=$)
- ☆ Natural join (\bowtie_*)
- ☆ Division operation (\div)
- ☆ Outer join
 - Left outer join ($\bowtie\lrcorner$)
 - Right outer join ($\lrcorner\bowtie$)
 - Full outer join ($\bowtie\lrcorner$)

1. Select (σ)

- ☆ Select tuples that satisfy a given predicate.
- ☆ It is denoted by lowercase Greek letter sigma (σ).
- ☆ Syntax: $\sigma_{\langle \text{selection_condition} \rangle}(\text{Relation})$.
- ☆ Example: $\sigma_{D_ID = 2}(\text{Employee})$.
- ☆ Comparison operators: $=$, \neq , $<$, \leq , $>$, and \geq .
- ☆ Connectives: AND (\wedge), OR (\vee) and NOT (\neg).



1. Select (σ)

Example 1: Write an RA expression to find all the instructors working in "Finance" department.

Solution:

$\sigma_{\text{Dept_Name} = \text{"Finance"}}(\text{INSTRUCTOR})$

Output:

ID	Name	Dept_Name	Salary
26589	Yusuf	Finance	95000
12547	Neil	Finance	80000

INSTRUCTOR			
ID	Name	Dept_Name	Salary
10101	John	Biology	65000
12121	Robin	Computer Science	90000
25252	Alya	Electrical	40000
26589	Yusuf	Finance	95000
54789	Ravi	Music	60000
78787	Raj	Physics	87000
87458	Jayant	History	75000
76985	Pratik	Computer Science	89000
12547	Neil	Finance	80000

1. Select (σ)

Example 3: Find all instructors who are working in "Finance" department and drawing the salary greater than \$87,000.

Solution:

$\sigma_{\text{Dept_Name} = \text{"Finance"} \wedge \text{Salary} > 87000} (\text{INSTRUCTOR})$

Output:

ID	Name	Dept_Name	Salary
26589	Yusuf	Finance	95000

INSTRUCTOR			
ID	Name	Dept_Name	Salary
10101	John	Biology	65000
12121	Robin	Computer Science	90000
25252	Alya	Electrical	40000
26589	Yusuf	Finance	95000
54789	Ravi	Music	60000
78787	Raj	Physics	87000
87458	Jayant	History	75000
76985	Pratik	Computer Science	89000
12547	Neil	Finance	80000

2. Project (Π)

- ☆ It returns its argument relation with **certain attributes left out**.
- ☆ It is a **unary** operator.
- ☆ It is denoted by the uppercase Greek letter **pi** (Π).
- ☆ Basically a relation is a **set**.
- ☆ In the result, the **duplicate rows** are eliminated.
- ☆ Syntax: $\Pi_{\text{Attribute1, Attribute2, ...}}(\text{Relation})$.

2. Project (Π)

Example 1: List all instructors' ID, name, and salary, but do not care about the dept_name.

Solution:

Π ID, Name, Salary (INSTRUCTOR)

INSTRUCTOR			
ID	Name	Dept_Name	Salary
10101	John	Biology	65000
12121	Robin	Computer Science	90000
25252	Alya	Electrical	40000
26589	Yusuf	Finance	95000
54789	Ravi	Music	60000
78787	Raj	Physics	87000
87458	Jayant	History	75000
76985	Pratik	Computer Science	89000
12547	Neil	Finance	80000

2. Project (Π)

Example 2: Find the name of all instructors in the Computer Science department.

Solution:

$\Pi_{\text{Name}} (\sigma_{\text{Dept_Name} = \text{"Computer Science"}} (\text{INSTRUCTOR}))$

Output:

Name
Robin
Pratik

INSTRUCTOR			
ID	Name	Dept_Name	Salary
10101	John	Biology	65000
12121	Robin	Computer Science	90000
25252	Alya	Electrical	40000
26589	Yusuf	Finance	95000
54789	Ravi	Music	60000
78787	Raj	Physics	87000
87458	Jayant	History	75000
76985	Pratik	Computer Science	89000
12547	Neil	Finance	80000

3. Union (U)

- ☆ Any relation is a **set**.
- ☆ Similar to **union** operation in Set Theory.
- ☆ It is a **binary** operator.
- ☆ It is a set of all objects that are a member of A, or B, or both.
- ☆ Like project, the **duplicate rows** are eliminated.
- ☆ It is denoted by U.
- ☆ Syntax: $\Pi_{\text{Column}}(\text{Relation}_1) \cup \Pi_{\text{Column}}(\text{Relation}_2)$



3. Union (U)

R	Alphabets
	A
	B
	C
	E
	F

S	Characters
	A
	B
	\$
	E
	F
	l
	#

R U S



R U S
A
B
C
E
F
\$
l
#

3. Union (U)

Example 1: List all customer names associated with the bank either as an account holder or a loan borrower.

Solution: $\Pi_{Cu_Name}(\text{Depositor}) \cup \Pi_{Cu_Name}(\text{Borrower})$

DEPOSITOR	Cu_Name	Acc_No
R	Tom	A-101
	Amy	A-502
	Rose	A-304
	John	A-689

BORROWER	Cu_Name	Loan_No
S	John	L-201
	Smith	L-658
	Rose	L-254
	Jack	L-547

Cu_Name
Tom
Amy
Rose
John
Smith
Jack

3. Union (U)

Example 2: Find the set of all courses taught in the Fall 2009 semester, the Spring 2010 semester, or both.

Solution:

To find the set of all courses taught in the Fall 2009 semester, we write:

$$\Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Fall"} \wedge \text{year} = 2009} (\text{SECTION}))$$

To find the set of all courses taught in the Spring 2010 semester, we write:

$$\Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Spring"} \wedge \text{year} = 2010} (\text{SECTION}))$$

To answer the query, we need the union of these two sets:

$$\Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Fall"} \wedge \text{year} = 2009} (\text{SECTION})) \cup \Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Spring"} \wedge \text{year} = 2010} (\text{SECTION}))$$

Two important conditions

For $R \cup S$ to be valid,

1. R and S must be of **same arity**.

2. For all i ,

Domain of the i^{th} attribute of R = **Domain** of i^{th} attribute of S .



4. Set Difference (-)

- ☆ It is like the same set difference in Set Theory.
- ☆ It is a binary operation.
- ☆ To find tuples that are in one relation but are not in another.
- ☆ $R - S = \text{Tuples in } R \text{ but not in } S.$
- ☆ It is denoted by minus (-) symbol.

4. Set Difference (-)

R	Alphabets
	A
	B
	C
	D
	E
	F

S	Characters
	A
	B
	\$
	L
	Y
	l
	#

R - S	R - S
	C
	D
	E
	F

4. Set Difference (-)

Example 1: List all customer names those who have a deposit account but not availed loan.

DEPOSITOR	Cu_Name	Acc_No	BORROWER	Cu_Name	Loan_No	R - S {	Cu_Name
R ↑	Tom	A-101	S ↑	John	L-201		Tom
	Amy	A-502		Smith	L-658		Amy
	Rose	A-304		Rose	L-254		
	John	A-689		Jack	L-547		

Solution: $\Pi_{Cu_Name} (Depositor) - \Pi_{Cu_Name} (Borrower)$

4. Set Difference (-)

Example 2: Find all the courses taught in the Fall 2009 semester but not in Spring 2010.

Solution:

To find the set of all courses taught in the Fall 2009 semester, we write:

$$\Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Fall"} \wedge \text{year} = 2009} (\text{SECTION}))$$

To find the set of all courses taught in the Spring 2010 semester, we write:

$$\Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Spring"} \wedge \text{year} = 2010} (\text{SECTION}))$$

To answer the query, we need the minus of these two sets:

$$\Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Fall"} \wedge \text{year} = 2009} (\text{SECTION})) - \Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Spring"} \wedge \text{year} = 2010} (\text{SECTION}))$$

Two important conditions

For $R - S$ to be valid,

1. R and S must be of same arity.
2. For all i ,

Domain of the i^{th} attribute of R = Domain of i^{th} attribute of S .

5. Cartesian Product (X)

- ☆ Cartesian product associates **every** tuple of R_1 with **every** tuple of R_2 .
- ☆ It is a **binary** operation.
- ☆ It is denoted by cross (X) symbol.
- ☆ $R_1 \times R_2 =$ All possible pairing.
- ☆ Same attribute may appear in R_1 and R_2 .
- ☆ $R = \text{Depositor} \times \text{Borrower}$.



5. Cartesian Product (X)

Example: Perform Depositor X Borrower.

DEPOSITOR	
Cu_Name	Acc_No
Tom	A-101
Rose	A-304

BORROWER	
Cu_Name	Loan_No
John	L-201
Smith	L-658
Rose	L-254
Jack	L-547

DEPOSITOR X BORROWER			
Cu_Name	Acc_No	Cu_Name	Loan_No
Tom	A-101	John	L-201
Tom	A-101	Smith	L-658
Tom	A-101	Rose	L-254
Tom	A-101	Jack	L-547
Rose	A-304	John	L-201
Rose	A-304	Smith	L-658
Rose	A-304	Rose	L-254
Rose	A-304	Jack	L-547

5. Cartesian Product (X)

- ☆ $R = \text{Depositor} \times \text{Borrower}.$
- ☆ $R = (\text{Depositor.Cu_Name}, \text{Depositor.Acc_No}, \text{Borrower.Cu_Name}, \text{Borrower.Loan_No}).$
- ☆ $R = (\text{Depositor.Cu_Name}, \text{Acc_No}, \text{Borrower.Cu_Name}, \text{Loan_No}).$



5. Cartesian Product (X)

Example : Find the names of all instructors in the Physics department together with the course id of all courses they taught.

Solution:

$\sigma_{\text{Dept_Name} = \text{"Physics"}} (\text{INSTRUCTOR } X \text{ TEACHES})$

$\sigma_{\text{Instructor.ID} = \text{Teaches.ID}} (\sigma_{\text{Dept_Name} = \text{"Physics"}} (\text{INSTRUCTOR } X \text{ TEACHES}))$

$\Pi_{\text{name, course_id}} (\sigma_{\text{Instructor.ID} = \text{Teaches.ID}} (\sigma_{\text{Dept_Name} = \text{"Physics"}} (\text{INSTRUCTOR } X \text{ TEACHES})))$

[or]

$\Pi_{\text{name, course_id}} (\sigma_{\text{Instructor.ID} = \text{Teaches.ID}} ((\sigma_{\text{Dept_Name} = \text{"Physics"}} (\text{INSTRUCTOR}) X \text{ TEACHES}))$

6. Rename (ρ)

- ☆ Relations in the database have names.
- ☆ The results of relational-algebra expressions **do not have a name**.
- ☆ It is useful to be able to give them names.
- ☆ It is a unary operation.
- ☆ It is denoted by the lowercase Greek letter **rho** (ρ).
- ☆ Syntax: $\rho_x (E)$



6. Rename (ρ)

- ☆ A relation r by itself is considered a **trivial** relational-algebra expression.
- ☆ Thus, the same rename operation can be applied to a relation r to get the same relation under a **new name**.
- ☆ The results of relational-algebra expressions **do not have a name**.
- ☆ $\rho_{x(A_1, A_2, \dots, A_n)}(\text{Expression})$

6. Rename (ρ)

Example : Find the names of all instructors in the Physics department together with the course id of all courses they taught and name the resultant relation as Ins_Phy.

Solution:

$\sigma_{\text{Dept_Name} = \text{"Physics"}} (\text{INSTRUCTOR X TEACHES})$

$\sigma_{\text{Instructor.ID} = \text{Teaches.ID}} (\sigma_{\text{Dept_Name} = \text{"Physics"}} (\text{INSTRUCTOR X TEACHES}))$

$\Pi_{\text{name, course_id}} (\sigma_{\text{Instructor.ID} = \text{Teaches.ID}} (\sigma_{\text{Dept_Name} = \text{"Physics"}} (\text{INSTRUCTOR X TEACHES})))$

[or]

$\Pi_{\text{name, course_id}} (\sigma_{\text{Instructor.ID} = \text{Teaches.ID}} ((\sigma_{\text{Dept_Name} = \text{"Physics"}} (\text{INSTRUCTOR}) \text{ X TEACHES}))$

$\rho_{\text{Ins_Phy}} (\Pi_{\text{name, course_id}} (\sigma_{\text{Instructor.ID} = \text{Teaches.ID}} ((\sigma_{\text{Dept_Name} = \text{"Physics"}} (\text{INSTRUCTOR}) \text{ X TEACHES}))))$

Set Intersection (\cap)

- ☆ Any relation is a **set**.
- ☆ Similar to **intersection** operation in Set Theory.
- ☆ It is a **binary** operator.
- ☆ It is a set of all objects that are a member of both A and B.
- ☆ It is denoted by \cap .
- ☆ Syntax: $\Pi_{\text{Column}}(\text{Relation_1}) \cap \Pi_{\text{Column}}(\text{Relation_2})$

Set Intersection (\cap)

R	Alphabets
	A 
	B
	C
	D
	E
	F

S	Characters
	A
	B
	\$
	E
	F
	l
	#

$R \cap S$

$R \cap S$
A
B
E
F

Set Intersection (\cap)

Example 1: Find the names of all customers who have deposited money and also availed loan.

Solution: $\Pi_{Cu_Name}(\text{Depositor}) \cap \Pi_{Cu_Name}(\text{Borrower})$

DEPOSITOR	Cu_Name	Acc_No
R	Tom	A-101
	Amy	A-502
	Rose	A-304
	John	A-689

BORROWER	Cu_Name	Loan_No
S	John	L-201
	Smith	L-658
	Rose	L-254
	Jack	L-547

$R \cap S$	Cu_Name
	Rose
	John

Set Intersection (\cap)

Example 2: Find the set of all courses taught in the Fall 2009 semester and the Spring 2010 semesters.

Solution:

To find the set of all courses taught in the Fall 2009 semester, we write:

$$\Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Fall"} \wedge \text{year} = 2009} (\text{SECTION}))$$

To find the set of all courses taught in the Spring 2010 semester, we write:

$$\Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Spring"} \wedge \text{year} = 2010} (\text{SECTION}))$$

To answer the query, we need the intersection of these two sets:

$$\Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Fall"} \wedge \text{year} = 2009} (\text{SECTION})) \cap \Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Spring"} \wedge \text{year} = 2010} (\text{SECTION}))$$

☆ $R \cap S = R - (R - S).$

Set Intersection (\cap): $R \cap S = R - (R - S)$

☆ $R \cap S = R - (R - S)$.

R	Alphabets
	A
	B
	C
	D
	E
	F

S	Characters
	A
	B
	\$
	E
	F
	l
	#

$R \cap S$
A
B
E
F

$R - S$
C
D

$R - (R - S)$
A
B
E
F

Assignment Operation (=)

- ☆ Similar to assignment operator in programming languages.
- ☆ Denoted by = or \leftarrow .
- ☆ It is a **binary** operator.
- ☆ The database might be **modified** if assignment to a permanent relation is made.
- ☆ Useful in the situation where it is required to write relational algebra expressions by using **temporary relation variable**.
- ☆ Example: $P = R \cap S$.
- ☆ Does not provide any additional power to relational algebra.

Assignment Operation (=)

- ☆ It provides a convenient way to express complex queries.
- ☆ Example:

Temp1 = Expression1

Temp2 = Expression2

Result = Temp1 ÷ Temp2

[or]

Temp1 ← R X S

Temp2 ← $\Pi_{\text{Course_id}} (\sigma_{\text{Semester} = \text{"Spring"} \wedge \text{year} = 2010} (\text{SECTION}))$

Result ← $\Pi_{R \cup S} (\text{Temp2})$

❖ Relational Algebra Operations – Binary

★ The JOIN Operation (\bowtie):

Syntax: $R \bowtie_{\langle \text{join condition} \rangle} S$

DEPARTMENT	DName	DNo	Mgr_SSN
	Research	2	553621425
	Finance	5	996856974

EMPLOYEE	SSN	FName	LName	DNo
	123658974	Alex	Smith	2
	553621425	Fred	Scott	5
	996856974	Elsa	David	5
	859689742	Peter	Williams	2

$\text{Dept_Mgr} \leftarrow \text{DEPARTMENT} \bowtie_{\text{Mgr_SSN} = \text{SSN}} \text{EMPLOYEE}$

❖ Relational Algebra Operations – Binary

★ The JOIN Operation (\bowtie):

DEPARTMENT	DName	DNo	Mgr_SSN
	Research	2	553621425
	Finance	5	996856974

EMPLOYEE	SSN	FName	LName	DNo
	123658974	Alex	Smith	2
	553621425	Fred	Scott	5
	996856974	Elsa	David	5
	859689742	Peter	Williams	2

Temp \leftarrow DEPARTMENT \bowtie EMPLOYEE

Dept_Mgr $\leftarrow \sigma_{(MgrSSN = SSN)}(Temp)$

❖ Relational Algebra Operations – Binary

★ The JOIN Operation (⋈):

DEPARTMENT	DName	DNo	Mgr_SSN
	Research	2	553621425
	Finance	5	996856974

EMPLOYEE	SSN	FName	LName	DNo
	123658974	Alex	Smith	2
	553621425	Fred	Scott	5
	996856974	Elsa	David	5
	859689742	Peter	Williams	2

Dept_Mgr	DName	DNo	Mgr_SSN	SSN	FName	LName	DNo
	Research	2	553621425	553621425	Fred	Scott	2
	Finance	5	996856974	996856974	Elsa	David	5

Result of the
JOIN
Operation



❖ Relational Algebra Operations – Binary

★ The THETA Join (θ):

<join condition> : $A_i \theta B_j$

Syntax: $R \bowtie_{A_i \theta B_j} S$

R	A_1	A_2
	20	25
	80	40

S	B_1
	50
	35

$A_i \Rightarrow$ Attribute of R

$B_j \Rightarrow$ Attribute of S

$\theta \Rightarrow \{ =, <, <=, >, >=, \neq \}$

$R \bowtie_{A_2 = B_1} S$

$R \bowtie_{A_2 < B_1} S$

$R \bowtie_{A_2 \leq B_1} S$

$R \bowtie_{A_2 > B_1} S$

$R \bowtie_{A_2 \geq B_1} S$

$R \bowtie_{A_2 \neq B_1} S$

❖ Relational Algebra Operations – Binary

★ The THETA Join (θ):

$R \bowtie_{A_2 > B_1} S$

R	A1	A2
	20	25
	80	40

S	B1
	50
	35

A1	A2	B1
20	25	50
20	25	35
80	40	50
80	40	35

Result of
"X"

A1	A2	B1
80	40	35

Result of " \bowtie "

❖ Relational Algebra Operations – Binary

★ The EQUIJOIN Operation:

- The only comparison operator used is “=”.

Dept_Mgr \leftarrow DEPARTMENT $\bowtie_{\text{Mgr_SSN} = \text{SSN}}$ EMPLOYEE

★ The NATURAL JOIN Operation (*):

- Can be performed only if there is a common attribute in between the relations.

❖ Relational Algebra Operations – Binary

★ The NATURAL JOIN Operation ($*$):

PROJECT	PID	PName	DNum
	101	ProjectX	1
	102	ProjectY	2
	103	ProjectZ	2

DEPARTMENT	DNo	Mgr_SSN
	1	553621425
	2	996856974

$\text{Proj_Dept} \leftarrow \text{PROJECT} * \rho_{(\text{DNum}, \text{Mgr_SSN})}(\text{DEPARTMENT})$

$\text{DEPT} \leftarrow \rho_{(\text{DNum}, \text{Mgr_SSN})}(\text{DEPARTMENT})$

$\text{Proj_Dept} \leftarrow \text{PROJECT} * \text{DEPT}$



❖ Relational Algebra Operations – Binary

★ The NATURAL JOIN Operation (*):

PROJECT	PID	PName	DNum
	101	ProjectX	1
	102	ProjectY	2
	103	ProjectZ	2

DEPARTMENT	DNNo	Mgr_SSN
	1	553621425
	2	996856974

Proj_Dept	PID	PName	DNum	Mgr_SSN
	101	ProjectX	1	553621425
	102	ProjectY	2	996856974
	103	ProjectZ	2	996856974

Result of " * "

❖ Relational Algebra Operations – Binary

★ The DIVISION Operation (\div):

Ex: To retrieve the Employee ID (**EID**) of the employees working on all projects.

EMPLOYEE	EID	PID
	1001	1
	1002	1
	1002	2
	1003	2

PROJECT	PID
	1
	2

Res	EID
	1002

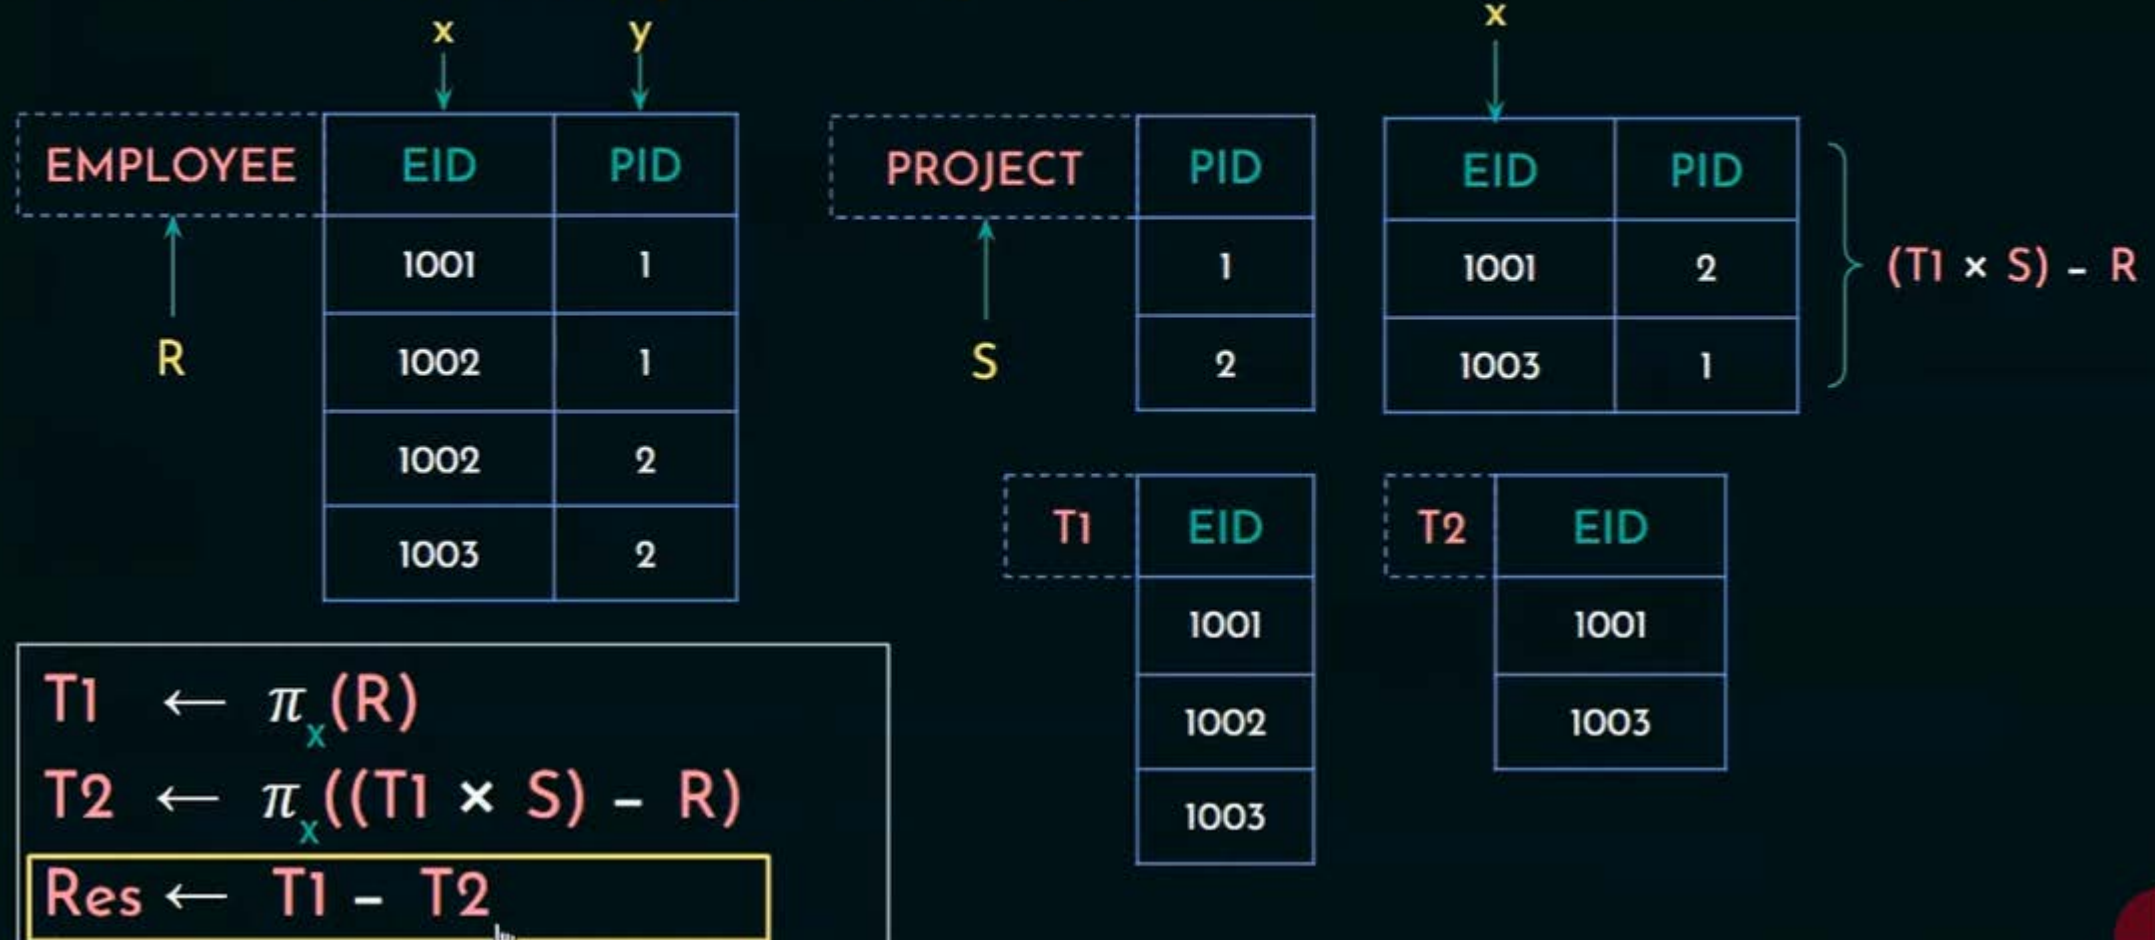
Result of the
" \div "
Operation

$Res \leftarrow EMPLOYEE \div PROJECT$




❖ Relational Algebra Operations – Binary

★ The DIVISION Operation (\div):




B-Tree (Dynamic Multilevel Index)



	Root	Intermediate Node
Max	p	p
Min	2	$\lceil p/2 \rceil$

B-Tree (Dynamic Multilevel Index)

10, 20, 30 Block Pointer
Tree Pointer



$$\text{Keys} = p - 1$$

$$\text{Rp} = p - 1$$


Order = p = Max. no of children = 4.

\downarrow
 B_p


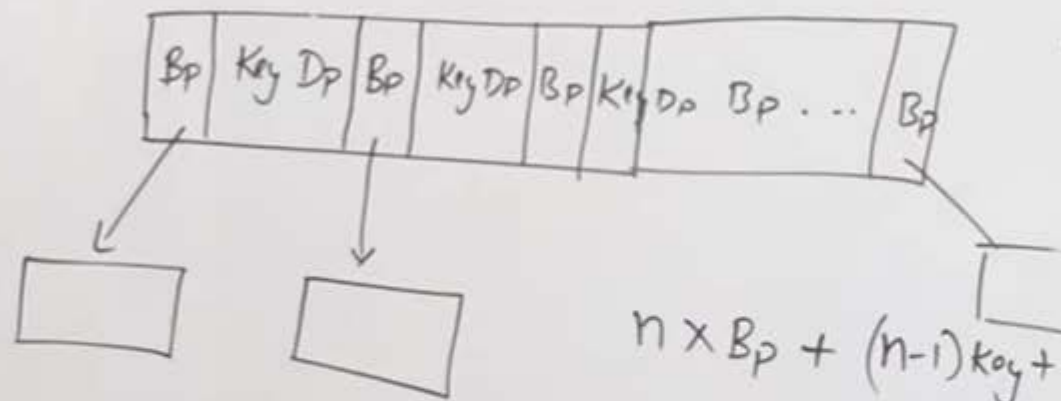
Children	Root	Intermediate Node
Max	p	p
Min	2	$\lceil p/2 \rceil$

Keys
Data Pointer

Record "



Consider a B-Tree with Key Size = 10 bytes, block size 512 bytes, data pointer is of size 8 bytes and block pointer is 5 bytes. Find the order of B-Tree?



$$n \times B_p + (n-1) \text{Key} + (n-1) D_p \leq \text{Block size}$$

$$n \times 5 + (n-1)(10+8) \leq 512$$

$$5n + 18n - 18 \leq 512$$

$$23n - 18 \leq 512$$

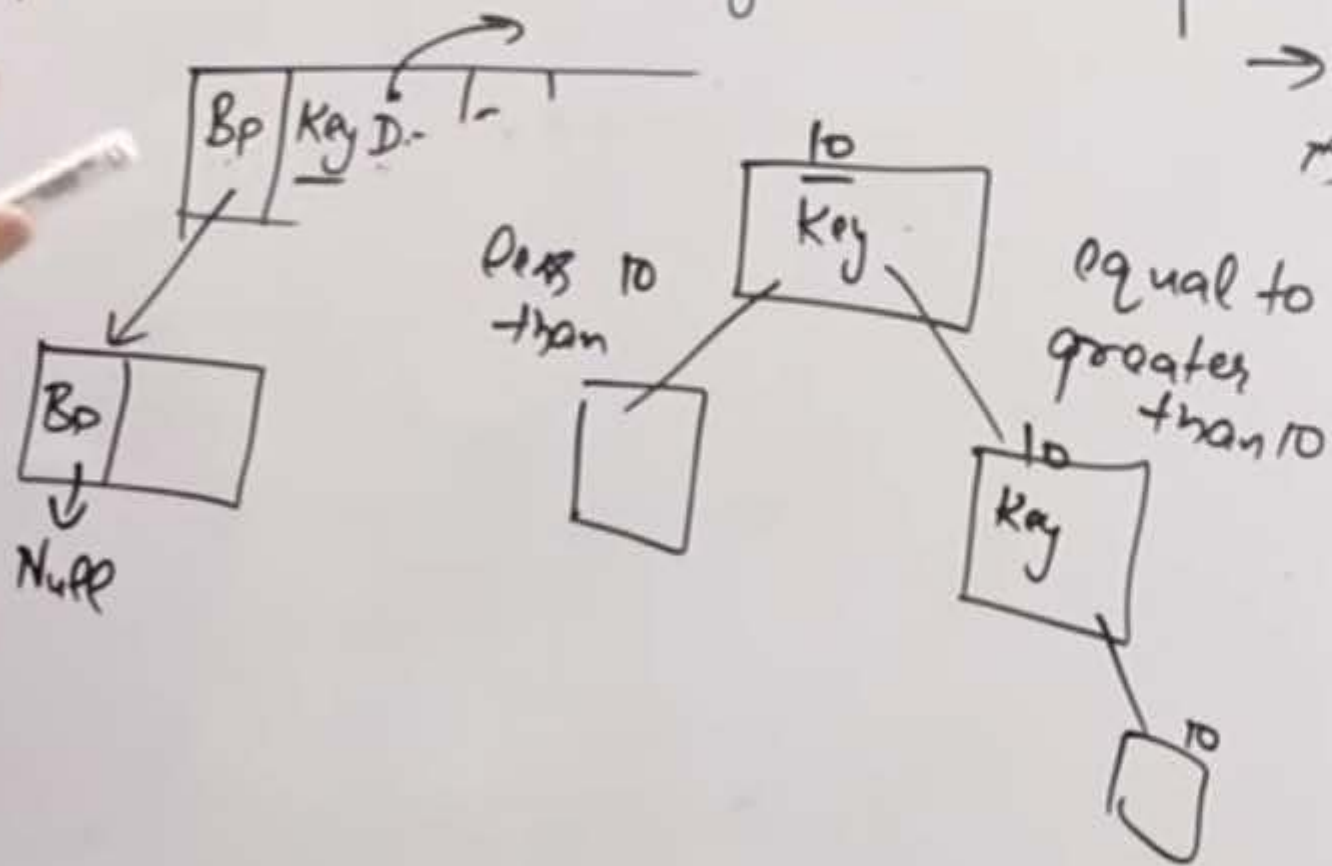
$$23n \leq 530$$

$$n \leq \frac{530}{23} = 23.04$$

23

B-Tree

- 1) Data is stored in leaf as well as internal nodes.
- 2) Searching is slower, deletion complex
- 3) No Redundant search key present
- 4) Leaf nodes not linked together




B⁺ Tree

- 1) Data is stored only in leaf nodes.
- 2) Searching is faster, deletion easy (directly from leaf node)
- 3) Redundant keys may present.
- 4) Linked together like linked list.





Internal node:



Consider a B⁺ Tree with Key Size = 10 bytes, block size = 512 bytes
 data pointer = 8 bytes and block pointer = 5 bytes. What is the order
 of leaf and non leaf node?

Non leaf.



$$n \times B_p + (n-1) \times \text{Key} \leq \text{Block size}$$

$$n \times 5 + (n-1) \times 10 \leq 512$$


$$5n + 10n - 10 \leq 512$$

$$15n \leq 522$$


$$n \leq \frac{522}{15} = 34.8$$

34

Key Dp pair.



Leaf:



$$n(\text{Key} + D_p) + B_p \leq \text{Block size}$$

$$n(10 + 8) + 5 \leq 512$$

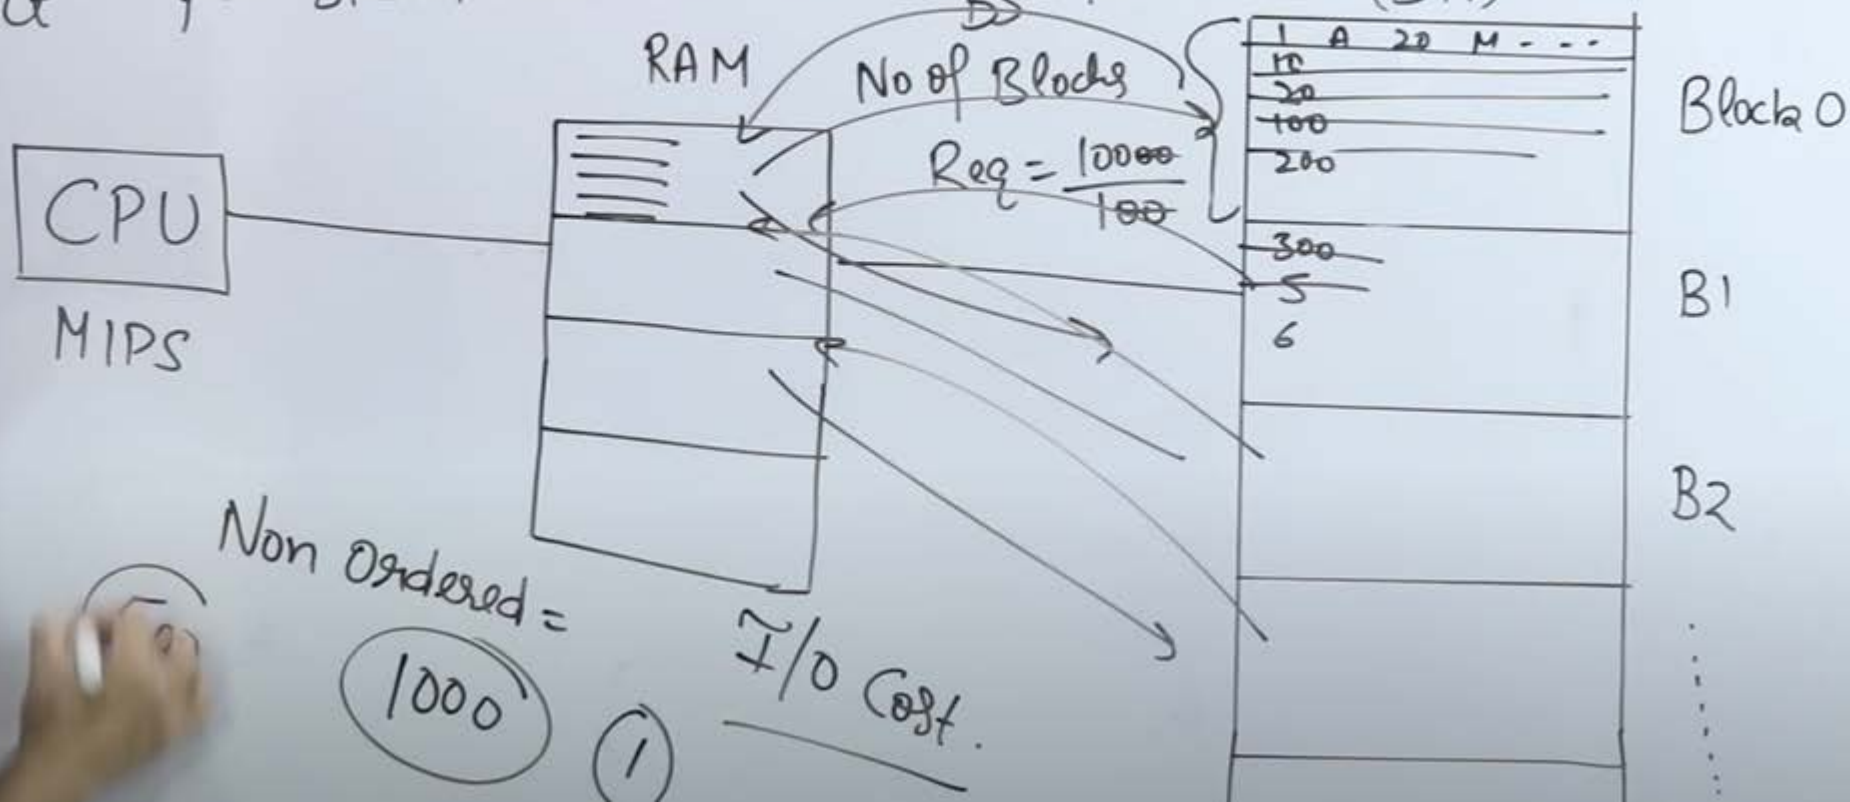
$$18n \leq 507$$

$$n \leq \frac{507}{18} = 28$$

28

Why Indexing is used? 10000 Records

Select * from Student where Rollno=1; BS = 100 Records Hard Disk (Slow)
(SM)





Without Indexing Consider a Hard Disk in which Block Size = 1000 Bytes, each record is of Size = 250 Bytes. If total no of records are 10000 and the data entered in Hard disk without any order (Unordered) (Ordered) what is avg time Complexity to search a record from HD?

Select *
from Student

Where $R_{all} = 345$

No of blocks Required = $\frac{10000}{4} = 2500$



1000 B $O(N)$

I/O Cost =
Best case = 1
Worst case = 2500 = N
Avg = $\frac{2500}{2} = 1250$

"Types of Indexes"

- 1) Primary
- 2) Clustered

Ordered
file

Unordered
file

Primary Index	Clustered Index
Secondary Index	Secondary Index

Secondary

10
4
6
12
2
1

1 10
2 9
7
6
...

Key
1
2
3
4
5
6
7
8
9
10
...

Non Key
1
2
3
4
5
6
7
8
9
10
...

2
5
2
3
9
8
5
2
9
1

Disk Architecture

Platters \rightarrow Surface \rightarrow Track \rightarrow Sectors \rightarrow Data



Actuator
Arm

$$8 \times 2 \times 256 \times 512 \times 512 \text{ KB}$$
$$2^3 \times 2^1 \times 2^8 \times 2^9 \times 2^9 \times 2^{10} \text{ B}$$
$$\text{Disk Size} = P \times S \times T \times S \times D$$

$$= 2^{40} \text{ B}$$


$$= \boxed{1 \text{ TB}}$$

$$\begin{array}{r} 0-15. \\ 0000 \\ 1111 \end{array} \Bigg] 16.$$


No. of bits required
to represent Disk Size.
40.

$$1 \text{ K} = 2^{10}$$
$$1 \text{ M} = 2^{20}$$
$$1 \text{ G} = 2^{30}$$
$$1 \text{ T} = 2^{40}$$


RAID 0



RAID 1



RAID 3



RAID 5

